

Training Small Language Models for Autoformalization to Prolog

Anonymous

Submitted for the Degree of Master of Science in
MSc Computer Science (Artificial Intelligence)



Department of Computer Science
Royal Holloway University of London
Egham, Surrey TW20 0EX, UK

August 29, 2026

Declaration

This report has been prepared on the basis of my own work. Where other published and unpublished source materials have been used, these have been acknowledged.

Word Count: 12,773

Student Name: Anonymous

Date of Submission:

Abstract

Large language models can produce convincing responses to reasoning tasks while still making logical errors. Autoformalization offers an alternative approach in which natural-language statements are translated into a formal representation that can be checked and executed by a symbolic reasoning system. This project investigates whether a small language model can be fine-tuned to translate natural-language reasoning statements into executable Prolog.

A controlled dataset of 2,000 natural-language and Prolog pairs was constructed across four levels of increasing complexity, covering facts, simple rules, multi-condition rules and executable reasoning problems. All target programs were automatically validated using SWI-Prolog. Qwen2.5-1.5B-Instruct was evaluated as a baseline and subsequently fine-tuned using QLoRA. Performance was measured using exact-match accuracy, syntactic validity, executability and reasoning correctness.

The baseline model achieved only 0.50% exact-match accuracy, with 39.70% of outputs being syntactically valid and executable. Following fine-tuning, exact-match accuracy increased to 81.41%, while syntactic validity and executability both reached 100%. Error analysis identified recurring weaknesses in complex relational and multi-condition structures. A targeted refinement dataset of 250 additional examples was subsequently constructed, after which exact-match accuracy increased to 87.94%, with 100% syntactic validity, executability and reasoning correctness on the 41 reasoning examples. However, because refinement was informed by errors observed on the test set, the final result is not a fully independent estimate of generalisation.

The results demonstrate that parameter-efficient fine-tuning can enable a small language model to learn reliable natural-language-to-Prolog translation within a controlled domain. They also show the value of symbolic execution for distinguishing syntactic correctness, executability and functional reasoning behaviour from exact textual agreement.

Contents

1	Introduction	1
2	Background and Related Work	4
2.1	Language Models and Logical Reasoning	4
2.2	Autoformalization	5
2.3	Prolog as a Symbolic Representation	6
2.4	Natural-Language Reasoning Datasets	7
2.5	Small Language Models and Parameter-Efficient Fine-Tuning	7
2.6	Position of This Project	8
3	Dataset Design and Construction	10
3.1	Dataset Design	10
3.2	Dataset Structure	12
3.3	Difficulty Levels	12
3.4	Vocabulary and Template Generation	13
3.5	Automated Prolog Validation	14
3.6	Training, Validation and Test Split	14
3.7	Final Dataset	15
4	Methodology and System Design	16
4.1	System Architecture	16
4.2	Model and Hardware	17
4.3	Input and Output Format	17
4.4	Baseline Evaluation	18
4.5	QLoRA Implementation	18
4.6	Refinement Training	19
4.7	Prolog Execution and Validation	19
4.8	Evaluation Metrics	20
4.9	Reproducibility and Experimental Control	21
5	Experiments and Results	23
5.1	Baseline Evaluation	23
5.2	Initial Fine-Tuned Model	24
5.3	Performance by Difficulty Level	26
5.4	Error Analysis	27
5.5	Targeted Refinement Dataset	28
5.6	Refined Model Results	29
5.7	Remaining Errors	30

5.8 Overall Comparison	32
6 Discussion	34
6.1 Effect of Fine-Tuning	34
6.2 Symbolic Execution as an Evaluation Tool	35
6.3 Reasoning Correctness	35
6.4 Remaining Difficulties	36
6.5 Effect of Targeted Refinement	37
6.6 Dataset and Experimental Limitations	37
7 Conclusion and Future Work	39
References	41

1 Introduction

Large language models (LLMs) can generate fluent natural language and perform tasks that appear to require reasoning. However, they can still produce incorrect or inconsistent answers, particularly when a task requires precise logical reasoning. A response may appear convincing while containing an invalid inference or reaching a conclusion that does not follow from the information provided. This limits the reliability of language models in tasks where correctness is more important than plausibility.

One approach to improving the reliability of language-model reasoning is autoformalization. Rather than asking a model to solve a reasoning problem entirely in natural language, autoformalization translates the problem into a formal representation that can be checked or executed by another system. This separates the interpretation of the natural-language problem from the execution of the resulting symbolic reasoning. If the generated representation can be executed by a symbolic reasoning system, its properties can be tested directly rather than assessed solely through the language model’s response.

This project investigates autoformalization using Prolog, a logic programming language based on facts, rules and queries [9]. Prolog provides a useful target representation because generated programs can be executed using an existing interpreter. A simple example of this translation is shown below.

Natural language: Every doctor is a professional.

Prolog: `professional(X) :- doctor(X).`

More complex statements can similarly be represented using multiple conditions, relations, inequality and negation. Once a natural-language reasoning problem has been converted into Prolog, queries can be executed against the generated program to determine whether the expected conclusions follow.

The project focuses specifically on small language models (SLMs) rather than substantially larger models. Smaller models require fewer computational resources and can be fine-tuned on consumer hardware, but their lower capacity raises the question of whether they can reliably learn the precise correspondence between natural-language statements and symbolic programs. This leads to the central research question: to what extent can a small language model be fine-tuned to translate natural-language reason-

ing problems into syntactically valid, executable and logically correct Prolog programs?

To investigate this question, a controlled dataset of natural-language and Prolog pairs was constructed. The dataset contains 2,000 examples divided across four levels of increasing complexity. These range from individual facts and simple rules to rules containing multiple conditions and examples requiring a Prolog query to be evaluated. The dataset was automatically validated using SWI-Prolog to ensure that the target programs were syntactically valid and executable. Reasoning examples were additionally checked to ensure that their queries produced the expected results. A group-aware training, validation and test split was used to reduce leakage between closely related examples.

The experiments use Qwen2.5-1.5B-Instruct as the small language model [8]. The unmodified model was first evaluated on the held-out test set to establish a baseline. Its performance demonstrated the difficulty of the task without task-specific training: only 0.5% of generated programs exactly matched the expected Prolog, while 39.7% were syntactically valid and executable. The model was subsequently fine-tuned on the constructed dataset using QLoRA. Following fine-tuning, exact-match accuracy increased to 81.41%, while all generated programs were syntactically valid and executable.

The remaining failures were examined through error analysis rather than treating aggregate accuracy as the only measure of performance. This analysis identified several recurring weaknesses, particularly in more complex multi-condition rules and specific relational structures. These observations were used to construct a targeted refinement dataset containing 250 additional examples focused on areas in which the first fine-tuned model performed poorly.

A second version of the model was then trained using this additional data and evaluated on the same test set. The final model achieved an exact-match accuracy of 87.94%, with 100% syntactic validity, 100% executability and correct results on all 41 reasoning examples in the test set. As the refinement data were informed by error patterns observed on this test set, these final results are treated as refinement results rather than as a fully independent estimate of generalisation.

These results also demonstrate why several evaluation criteria are useful for autoformalization. Exact string matching provides a strict measure of whether a generated program is identical to the reference representation, but it does not by itself indicate whether the generated output is usable. Syntax validation determines whether the output forms valid Prolog, while

execution determines whether it can be successfully processed by the Prolog interpreter. Reasoning evaluation provides an additional functional measure by testing whether generated programs produce the expected answer to a query. This distinction is important because a generated symbolic representation can differ from a reference string while still retaining useful executable behaviour.

The main contribution of this project is an end-to-end investigation of small-model autoformalization to Prolog. It includes the construction and validation of an NL-Prolog dataset, the establishment of an untuned baseline, parameter-efficient fine-tuning of a 1.5-billion-parameter language model, execution-based evaluation using SWI-Prolog, systematic analysis of translation errors and targeted refinement based on the observed failures. The experiments show that task-specific fine-tuning can substantially improve a small language model’s ability to generate valid executable Prolog, while also identifying translation structures for which exact generation remains more difficult.

The remainder of this dissertation develops this investigation in detail:

- Chapter 2 introduces the relevant background and related work in language-model reasoning, autoformalization, Prolog and parameter-efficient fine-tuning.
- Chapter 3 describes the construction, validation and splitting of the NL-Prolog dataset.
- Chapter 4 presents the model, training procedure, system design and evaluation methodology.
- Chapter 5 reports the baseline, fine-tuning, error analysis and refinement experiments.
- Chapter 6 discusses the significance and limitations of the results, including the distinction between exact-match and execution-based evaluation.
- Chapter 7 summarises the findings of the project and identifies directions for future work.

2 Background and Related Work

This chapter introduces the concepts and previous research most relevant to the project. The project lies at the intersection of language models, symbolic reasoning and autoformalization. This central approach uses a language model to translate natural-language statements into Prolog rather than requiring the model to perform the entire reasoning process directly in natural language. This section therefore considers language-model reasoning, neuro-symbolic approaches, Prolog as a target formal language, existing reasoning datasets, and parameter-efficient methods for adapting smaller language models.

2.1 Language Models and Logical Reasoning

Transformer-based language models learn statistical relationships between tokens from large collections of text. This enables modern models to perform a wide range of tasks through natural-language prompting, including question answering, code generation and tasks that appear to require multi-step reasoning. However, generating a plausible sequence of tokens is not equivalent to following a formally defined reasoning procedure. This distinction becomes important when the correctness of each inference matters.

Research into logical reasoning with language models predates many current large language models. Clark et al. [2] introduced RuleTaker, in which transformer models were trained to answer questions from natural-language facts and rules. This demonstrated that transformer models could achieve high accuracy on synthetically generated logical reasoning problems and could generalise to some problems requiring deeper inference chains than those encountered during training. The authors described these models as “soft theorem provers”, since the reasoning process was learned from examples rather than performed through an explicit symbolic inference system.

RuleTaker demonstrates that transformer models can learn patterns corresponding to deductive reasoning. However, this distinction between learning the behaviour of a reasoning process and executing a formal reasoning procedure remains important. A neural model may produce the correct answer without providing a representation that can independently be checked by a deterministic solver. This motivates approaches in which language understanding and formal inference are separated.

A related line of research has therefore investigated neuro-symbolic reasoning, combining the flexibility of neural language models with the explicit inference procedures of symbolic systems. This approach does not require

the language model to both interpret a problem and solve it; instead, the language model can act as a translator or semantic parser. The resulting formal representation is then passed to a symbolic system responsible for inference.

2.2 Autoformalization

Autoformalization is the automatic translation of informal natural-language statements into a formal representation. This was investigated directly by Wu et al. [10], who used large language models to translate mathematical statements into Isabelle/HOL. Their experiments showed that language models could correctly formalize a portion of mathematical competition problems and that the generated formalizations could subsequently support automated theorem proving. This demonstrated that language models could serve as an interface between informal human language and representations suitable for formal verification.

The importance of autoformalization is that formal output can be subjected to checks that are unavailable for ordinary natural-language responses. This means that a natural-language answer may be grammatically correct and persuasive while still containing an invalid inference. A formal representation instead has a defined syntax and semantics and can be processed by an appropriate interpreter, theorem prover or solver. This translation may still be incorrect, but errors can potentially be detected through parsing, execution or comparison with expected behaviour.

Pan et al. [7] developed this idea further with Logic-LM, a framework in which an LLM first translates a natural-language reasoning problem into a symbolic formulation. A deterministic symbolic solver then performs the actual inference. Logic-LM was evaluated on five reasoning datasets, including ProofWriter and FOLIO, and produced substantial improvements over both standard prompting and chain-of-thought prompting. This framework also included a refinement stage that used error messages from the symbolic solver to revise incorrect formalizations.

A similar approach is used by Olausson et al. [6] in LINC, or Logical Inference via Neurosymbolic Computation. This system treats the language model as a semantic parser that translates natural-language premises and conclusions into first-order logic. Deductive reasoning is then delegated to an external theorem prover. Their experiments on ProofWriter and FOLIO found that this separation could outperform direct chain-of-thought reasoning. This approach also produced strong results when using an open-source model smaller than the largest proprietary models evaluated in the study.

These systems support the central principle behind the present project: a language model does not necessarily need to perform every part of logical reasoning internally. This allows it to perform the task for which natural-language models are particularly useful, interpreting language, while a symbolic system handles formally defined inference.

2.3 Prolog as a Symbolic Representation

This project uses Prolog as the target language for autoformalization. Prolog is a logic programming language in which programs are primarily expressed using facts and rules. This knowledge base can then be queried to determine whether particular conclusions follow from those facts and rules.

For example, the statements:

```
doctor(mary).
professional(X) :- doctor(X).
```

represent the fact that Mary is a doctor and the rule that every doctor is a professional. A query such as:

```
professional(mary).
```

can then be evaluated by the Prolog system. Prolog uses mechanisms including unification and proof search to determine whether a query follows from the available knowledge base.

This makes Prolog useful for the current project in two ways. First, its relatively compact syntax provides a structured target for natural-language translation. This allows natural-language statements describing properties and relations to be represented as predicates, while conditional statements can be represented as rules. Second, generated programs can be executed. This makes it possible to distinguish several forms of correctness: whether the generated text exactly matches a reference program, whether it is syntactically valid Prolog, whether it executes successfully, and whether executing a query produces the expected result.

Recent neuro-symbolic work has also explored Prolog specifically. Yang et al. [11] use Prolog interpreters as part of a neuro-symbolic reasoning framework and argue that, once a problem has been correctly translated into Prolog, the interpreter can provide reliable symbolic reasoning and expose its intermediate search process. More recently, Mellgren et al. [5] investigated training Qwen2.5-3B-Instruct to use Prolog as an external reasoning tool. This work considered execution, syntax and semantic correctness as training signals and highlighted a distinction between obtaining a correct final answer and producing a genuinely auditable symbolic program.

This distinction is relevant to the evaluation performed in this project. A model that produces valid Prolog has achieved something different from a model that merely produces the expected final answer. This executable symbolic representation creates an intermediate artefact that can itself be inspected and evaluated.

2.4 Natural-Language Reasoning Datasets

Training a model to translate natural language into a formal representation requires paired examples containing linguistic statements and their corresponding logical structures. This means that existing logical reasoning datasets provide useful examples of how controlled reasoning problems can be constructed.

RuleTaker uses synthetic natural-language theories containing facts and rules, with questions requiring models to determine whether conclusions follow from them. This allows reasoning depth and logical structure to be controlled systematically. Bao et al. [1] extended this direction through PARARULE-Plus, a large dataset designed for multi-step deductive reasoning over natural language. This dataset contains reasoning chains at depths from two to five and approximately 400,000 examples across its different depth levels. It uses entities such as people and animals together with relationships and attributes and follows a closed-world reasoning setting.

These datasets influenced the design space of the current project, particularly the use of controlled natural-language templates and progressively more complex reasoning structures. However, this objective differs from ordinary question answering over these datasets. The desired model output is not simply a true or false label. Instead, the model must generate the underlying Prolog representation itself.

For this reason, a dedicated NL-Prolog dataset was constructed rather than directly training on PARARULE-Plus. This provides precise control over the expected Prolog syntax, predicate structures and difficulty levels and allows every target program to be automatically checked with the same Prolog implementation used during evaluation. This dataset construction is described in Chapter 3.

2.5 Small Language Models and Parameter-Efficient Fine-Tuning

Many autoformalization and neuro-symbolic studies use comparatively large language models. This project instead investigates whether a much smaller

model can acquire the required translation behaviour through targeted fine-tuning.

The selected model is Qwen2.5-1.5B-Instruct, part of the Qwen2.5 family. Its technical report describes a family of models offered at several parameter scales and improved through both pre-training and post-training. The instruction-tuned variants are designed to follow user instructions and generate structured responses, making the 1.5-billion-parameter version a suitable starting point for a constrained generation task.

Fine-tuning all parameters of a language model can require substantial memory and computational resources. This requirement is reduced by Low-Rank Adaptation (LoRA), introduced by Hu et al. [4], which freezes the original model weights and introduces a smaller number of trainable low-rank matrices. Only these additional parameters need to be optimised for the downstream task. This substantially reduces the number of trainable parameters while retaining much of the benefit of task-specific adaptation.

QLoRA, introduced by Dettmers et al. [3], extends this approach by combining LoRA with quantisation. This method stores the pretrained model using 4-bit quantised weights while gradients are propagated into the trainable LoRA adapters. Dettmers et al. demonstrated that this approach can greatly reduce memory requirements while maintaining competitive fine-tuning performance.

This makes QLoRA particularly appropriate to the aims of this project because it makes fine-tuning feasible on limited hardware. The investigation can therefore focus on whether a small model can learn NL-Prolog translation without requiring the computational resources associated with full fine-tuning or substantially larger language models.

2.6 Position of This Project

Previous research provides strong evidence for combining neural language understanding with symbolic reasoning. Autoformalization work has shown that language models can translate informal statements into formal languages, while Logic-LM and LINC demonstrate the value of delegating inference to deterministic symbolic solvers. RuleTaker and PARARULE-Plus further show how controlled natural-language datasets can be used to investigate logical reasoning systematically. Recent Prolog-based work also demonstrates growing interest in using executable logic programs as an interface between language models and formal reasoning.

The present project focuses on a narrower question within this area. This project does not test whether a large general-purpose model can

be prompted to produce symbolic representations; instead, it investigates whether a 1.5-billion-parameter model can be explicitly trained to perform NL-Prolog translation. It also evaluates the generated output at several levels, separating exact textual agreement from syntax validity, executability and reasoning correctness. This provides a practical investigation of whether relatively small language models can learn a reliable interface between natural language and executable symbolic reasoning.

3 Dataset Design and Construction

The dataset is a central component of this project because fine-tuning requires paired examples showing the relationship between natural-language statements and their expected Prolog representations. Existing datasets such as RuleTaker, ProofWriter and PARARULE-Plus provide large collections of controlled natural-language reasoning problems. However, their primary objective is usually to determine whether a conclusion follows from a collection of facts and rules rather than to train a model to generate executable Prolog directly. A dedicated dataset was therefore constructed for this project.

The final dataset contains 2,000 natural-language and Prolog examples divided into four levels of increasing complexity. The design aimed to provide a controlled progression from basic translation to more complex rules and executable reasoning. Each example was generated together with meta-data describing its level, category and template group. The complete dataset was validated using SWI-Prolog before being divided into training, validation and test sets.

3.1 Dataset Design

Synthetic generation was chosen because it provides direct control over both sides of each training pair. Similar approaches have previously been used for logical reasoning datasets. RuleTaker generates theories containing natural-language facts and rules from a configurable grammar and derives labels using a reasoning engine. ProofWriter extends this style of dataset with rule bases, questions and proofs of different depths. PARARULE-Plus similarly generates large numbers of natural-language reasoning examples involving people, animals, attributes and relationships.

Together, these approaches demonstrate the usefulness of controlled synthetic data when the underlying logical structure of every example must be known. The same principle was used in this project, although the target differs. Rather than generating a natural-language theory followed by a classification label, each example pairs natural language directly with its expected Prolog program.

An example of a simple pair is:

Natural language: Mary is a doctor.
Prolog: `doctor(mary)`.

A rule can be represented as:

Natural language: Every doctor is a professional.
Prolog: professional(X) :- doctor(X).

The direct pairing makes the required model behaviour explicit. The model receives a natural-language statement and is trained to return only its Prolog representation.

The dataset was generated programmatically rather than written manually. Templates defined corresponding natural-language and Prolog structures, while entities and predicates were varied to produce individual examples. This reduced the amount of manual annotation required and ensured that the intended symbolic representation was known at generation time.

The overall generation procedure is summarised in Algorithm 1. Examples were generated according to the requirements of each difficulty level, with predefined templates controlling the relationship between the natural-language statement and its corresponding Prolog representation. Reasoning examples additionally included a query and an expected result.

Algorithm 1 NL-Prolog Dataset Generation

Require: Level specifications, NL templates, Prolog templates, entity names and predicates

Ensure: Dataset of NL-Prolog examples

```

1: Initialise empty dataset
2: for each difficulty level do
3:   Load the templates associated with the level
4:   while target number of examples has not been reached do
5:     Select a template
6:     Select the required entities and predicates
7:     Generate the natural-language statement
8:     Generate the corresponding Prolog representation
9:     Assign the level, category and template group
10:    if the example requires reasoning then
11:      Generate the associated query
12:      Assign the expected query result
13:    end if
14:    if the generated example is unique then
15:      Add the example to the dataset
16:    end if
17:  end while
18: end for
19: Save the completed dataset

```

However, purely template-based generation introduces the risk that a model learns surface patterns rather than the underlying translation. To reduce this effect, multiple natural-language forms, entities and predicate combinations were included. Examples were also assigned to template groups so that closely related examples could later be kept within the same dataset split.

3.2 Dataset Structure

Each dataset entry was stored in JSONL format. The principal fields were `id`, `level`, `category`, `natural_language`, `prolog`, `template_group` and `nl_template`. Reasoning examples additionally contain a `query` and an `expected_result`.

The `id` provides a unique identifier that allows individual examples to be traced during validation and error analysis. The `level` records the intended difficulty of the example, while `category` describes the type of translation. `natural_language` contains the model input and `prolog` contains the expected output.

The `template_group` field identifies examples derived from the same underlying logical pattern. This became important during splitting because two sentences can use different names or surface wording while representing almost the same translation task. The `nl_template` field provides additional information about the linguistic form from which the example was generated.

Reasoning examples contain two further fields. The `query` specifies the Prolog goal that should be executed after loading the generated program, while `expected_result` records whether the query should succeed. These fields allow the dataset to test functional reasoning behaviour in addition to translation.

JSONL was used because each example can be processed independently while retaining structured metadata. It also allowed the same records to be reused throughout generation, validation, splitting, training and evaluation.

3.3 Difficulty Levels

The 2,000 examples were divided into four levels. Each level introduces additional structural complexity.

Level 1 contains 500 facts. These examples establish the simplest correspondence between natural-language statements and Prolog predicates. Both unary predicates, such as properties of a single entity, and binary predicates describing relationships between two entities were included.

Level 2 contains 600 simple rules. These examples introduce variables and implication. A natural-language statement such as “Every doctor is a professional” must be converted into a Prolog rule with the conclusion before `:-` and its condition afterwards. This requires more structural transformation than copying information into a fact.

Level 3 contains 500 multi-condition rules. These introduce conjunctions and more complex logical structures, including inequality and negation-as-failure. For example, a rule may require several conditions to be satisfied before its conclusion follows. The level was designed to test whether the model could preserve multiple relationships and variables while producing valid Prolog syntax.

Level 4 contains 400 reasoning examples. These contain sufficient facts and rules for a specified query to be executed. The expected result of the query is stored alongside each example. Evaluation can therefore move beyond checking whether the generated output resembles the reference representation and test whether the resulting program supports the expected inference.

The use of progressively more difficult structures follows a similar principle to datasets such as ProofWriter and PARARULE-Plus, where reasoning problems are organised according to inference depth. The levels in this project are not equivalent to their reasoning depths because they measure translation complexity rather than only the number of inference steps. Nevertheless, both approaches provide a controlled way to examine where model performance begins to degrade.

3.4 Vocabulary and Template Generation

A controlled vocabulary was used when generating examples. Twenty-six entity names were included, such as `ash`, `mary`, `olivia`, `ray`, `violet` and `zelda`. Predicates were divided into unary and binary forms. Unary predicates represented properties or classes, while binary predicates represented relationships such as `calls/2`, `admires/2`, `parent/2`, `likes/2`, `teacher_of/2` and `student_of/2`.

The controlled vocabulary was intended to keep the emphasis on learning the mapping between linguistic and logical structure. RuleTaker and PARARULE-Plus use a related strategy, combining limited collections of entities, attributes and relationships to create large numbers of controlled reasoning problems.

Templates were created for each category and instantiated with different entities and predicates. More complex templates explicitly controlled

variable placement because reversing the arguments of a binary predicate can change the meaning of the generated program. This was particularly important for relationships such as `parent/2` and `teacher_of/2`.

Several issues were discovered during dataset development. These included incorrect article selection such as “a” versus “an”, an incorrect predicate mapping, duplicated Level 3 structures, and errors in the expected outcomes of some reasoning examples. The problems were corrected before the dataset was frozen. Iterative checking was important because systematic errors in automatically generated targets would otherwise teach the model incorrect translations.

3.5 Automated Prolog Validation

Automatic validation was included as part of the dataset construction pipeline. Each generated target was written to a temporary Prolog file and checked using SWI-Prolog. Examples were tested for both syntactic validity and successful loading by the interpreter.

Reasoning examples underwent an additional check. After loading the facts and rules, the associated query was executed and its result was compared with `expected_result`. This ensured that the expected answer stored in the dataset agreed with the behaviour of the Prolog program.

The use of a symbolic engine to derive or verify dataset labels has precedent in related work. RuleTaker’s generator, for example, can pass generated theories through a reasoning engine to obtain gold labels. The approach used here applies the same general principle directly to the target representation: because Prolog is executable, the generated target programs can be checked before they are used for model training.

Validation initially revealed errors in the generated data, particularly in the reasoning examples. The errors were corrected and the validation process was repeated until all 2,000 examples were syntactically valid and executable, and all 400 reasoning examples produced their expected results. The validated dataset was then frozen as the first dataset version so that subsequent experiments would operate on a consistent collection of examples.

3.6 Training, Validation and Test Split

The validated dataset was divided into training, validation and test sets using an approximately 80/10/10 split. This produced 1,597 training examples, 204 validation examples and 199 test examples. A fixed random seed

of 4242564 was used so that the split could be reproduced.

A purely random split at the level of individual examples was avoided. Template-generated datasets can contain examples that differ only in entity names or minor wording. Placing closely related examples in both training and test data would make evaluation less informative because the model could encounter almost the same pattern during training.

Examples were therefore grouped using `template_group`, and groups rather than individual examples were assigned to dataset splits. This reduces direct leakage between structurally related examples and provides a more meaningful test of whether the model has learned mappings that extend beyond individual generated examples.

The test set was kept fixed throughout the baseline, initial fine-tuning and refinement experiments. Maintaining the same test data made direct comparisons between model versions more reliable.

3.7 Final Dataset

The resulting dataset contains 2,000 validated NL-Prolog pairs covering facts, simple rules, multi-condition rules and executable reasoning problems. Its construction combines ideas from controlled synthetic reasoning datasets with direct symbolic program generation. Unlike datasets designed primarily for true-or-false reasoning classification, each example provides the formal representation itself as the target.

The dataset serves two roles in the project. First, it provides supervised examples for teaching a small language model how natural-language logical structures correspond to Prolog. Second, its executable targets provide the basis for evaluation using syntax checking, execution and reasoning behaviour. This connection between dataset construction and executable evaluation is central to the experimental methodology described in the following chapter.

4 Methodology and System Design

This chapter describes the implementation used to investigate whether a small language model can translate natural-language reasoning statements into executable Prolog. The overall system follows the pipeline shown below.

System Pipeline

Natural-Language Input → Language Model → Generated Prolog → SWI-Prolog Validation and Execution → Evaluation

The language model generates a Prolog representation from the natural-language input. The generated program is then passed to SWI-Prolog, which is used to validate and execute the output. The resulting predictions and execution results are subsequently used to calculate the evaluation metrics described later in this chapter.

The same evaluation pipeline was used for the baseline, initial fine-tuned model and refined model. All three models were evaluated using the same test set and evaluation procedure. However, the refined model was developed after analysing error patterns observed on this test set, so its final evaluation is not fully independent of the refinement process.

4.1 System Architecture

Each dataset example contains a natural-language input and, depending on its level, a reference Prolog program, a query and an expected result. During inference, the natural-language input is passed to the model with an instruction to return only Prolog code. The generated response is then:

1. stored for later analysis;
2. compared with the reference Prolog;
3. written to a temporary Prolog file;
4. checked and executed using SWI-Prolog; and
5. evaluated using the metrics described in Section 4.8.

This implementation separates translation errors from execution errors. A response may fail exact match while still being valid Prolog, or it may

match the intended structure but fail because of a syntax error. For reasoning examples, a program can also be syntactically valid but produce an incorrect answer to the associated query. The evaluation pipeline therefore records these outcomes separately.

4.2 Model and Hardware

The model used in all experiments was **Qwen2.5-1.5B-Instruct** [8]. It was selected because its instruction-tuned format is compatible with the task: the model receives an instruction and a natural-language statement, then generates a formal response.

The 1.5B parameter version was chosen to keep training and inference practical on the available hardware. Experiments were conducted using an NVIDIA GeForce RTX 4070 with 12 GB of GPU memory. This constraint made full-parameter fine-tuning impractical and motivated the use of QLoRA.

The original model was used for the baseline evaluation. Fine-tuned experiments used the same base model with separately saved LoRA adapters, allowing the baseline and adapted models to be compared without modifying the original model weights.

4.3 Input and Output Format

The task was implemented as supervised instruction tuning. Each training example was converted into an instruction-response pair. The input contained the natural-language statement, while the response contained the reference Prolog program.

The system instruction required the model to return only the generated Prolog without additional explanatory text. An example of the input format and expected response is shown below.

System instruction:

Return only the Prolog code. Do not include explanations or Markdown code fences.

User input:

Every doctor is a professional.

Expected response:

professional(X) :- doctor(X).

This instruction was used during both baseline and fine-tuned inference. It was intended to prevent explanatory text from being included in the generated response, since such text would interfere with exact-match evaluation and Prolog execution.

The same formatting procedure was applied to the training, validation and test examples. Predictions were saved in JSONL format together with the input, reference output, generated output and evaluation results.

4.4 Baseline Evaluation

Before fine-tuning, Qwen2.5-1.5B-Instruct was evaluated on the held-out test set. The baseline model received the fixed output instruction and one natural-language test example at a time. No training examples were included in the prompt.

The baseline responses were saved separately from later experiments. The test set contained 199 examples and remained unchanged throughout the project. This provided a fixed reference point for measuring the effect of QLoRA fine-tuning and subsequent refinement.

The baseline evaluation used the same output normalisation, Prolog validation and reasoning checks as the fine-tuned models. Consequently, differences between the models reflect changes in model parameters rather than changes in the evaluation procedure.

4.5 QLoRA Implementation

QLoRA was used to adapt Qwen2.5-1.5B-Instruct to the NL-Prolog translation task. The implementation combined 4-bit loading of the base model with trainable LoRA adapter layers.

The base model was loaded in 4-bit quantised form. Its original parameters were frozen during training. Trainable low-rank adapter parameters were inserted into selected model layers, and only these adapter parameters were updated. The resulting training process required substantially less GPU memory than updating all model parameters.

The first fine-tuning run used the 1,597 examples in the training split. The 204 validation examples were kept separate and were not included in the training data. Training was performed for three epochs with a per-device batch size of two. The random seed was set to 4242564.

Training was carried out using the TRL supervised fine-tuning pipeline. The model used float16 computation where supported, while the base model remained loaded using 4-bit quantisation. The trained adapter was saved

separately from the base Qwen model. During inference, the adapter was loaded together with the original base model.

The implementation therefore did not create a new full copy of Qwen2.5-1.5B. Instead, it retained the quantised base model and stored the learned task-specific changes in the adapter. This made it possible to maintain separate versions for the initial fine-tuned model and the refined model while using the same base model.

QLoRA was selected because it allowed the experiment to be performed within the memory limits of the RTX 4070. The purpose of the experiment was to evaluate task adaptation, not to compare different quantisation methods. Quantisation was therefore used as an implementation strategy to make fine-tuning feasible.

4.6 Refinement Training

After evaluating the first fine-tuned model, generated outputs were grouped by difficulty level, category and template type. Incorrect and non-exact outputs were inspected to identify recurring translation problems.

The analysis showed that errors were concentrated in particular structures. Multi-condition Level 3 examples were more difficult than simple facts and rules. Errors also occurred repeatedly in examples involving inequality, negation and binary predicates.

A refinement dataset containing 250 additional examples was then created. These examples placed greater emphasis on the structures associated with the observed errors. The refinement data were combined with the original training data rather than replacing them.

A second QLoRA training run was performed using this combined dataset. The resulting adapter was saved separately from the first adapter. The validation and test sets were not regenerated or expanded. This ensured that the refined model could be compared directly with the baseline and first fine-tuned model using the same evaluation examples.

4.7 Prolog Execution and Validation

Generated responses were evaluated using **SWI-Prolog** [9]. For each prediction, the generated text was written to a temporary `.pl` file. The file was passed to the SWI-Prolog executable, and the subprocess return code and diagnostic output were recorded. The temporary file was then deleted.

The evaluation distinguished between the following outcomes:

- **Syntax validity:** whether SWI-Prolog could parse the generated program;
- **Executability:** whether the generated program could be loaded and processed successfully; and
- **Reasoning correctness:** whether the program produced the expected result for the associated query.

For Level 4 examples, the query stored in the dataset was executed against the generated knowledge base. The result was compared with the expected Boolean value. A program could therefore be syntactically valid but receive an incorrect reasoning result if the generated facts or rules did not represent the original natural-language statement correctly.

The generated Prolog was also compared with the reference program before or alongside execution. This allowed textual translation accuracy to be reported independently from symbolic execution performance.

4.8 Evaluation Metrics

Four metrics were used.

Exact match measures whether the generated Prolog was identical to the reference after the normalisation applied by the evaluation script. If N is the number of examples and M is the number of exact matches:

$$\text{Exact Match} = \frac{M}{N}$$

Exact match is a strict translation measure. A non-exact response may still be valid Prolog, so exact match was not used as the only measure of performance.

Syntax validity is the proportion of generated programs accepted by SWI-Prolog without a parsing error.

Executability is the proportion of generated programs that could be loaded and processed successfully by SWI-Prolog.

Reasoning accuracy applies to the Level 4 examples. If R_c is the number of examples for which the generated program produced the expected query result and R is the total number of reasoning examples:

$$\text{Reasoning Accuracy} = \frac{R_c}{R}$$

These metrics were calculated for the baseline, first fine-tuned model and refined model. Reporting them separately makes it possible to determine

whether a model improved primarily in textual translation, Prolog validity or logical reasoning.

4.9 Reproducibility and Experimental Control

The dataset split was fixed before model comparison. The same 1,597 training examples, 204 validation examples and 199 test examples were used throughout the main experiments.

To support reproducibility, the main training, inference and software configuration used throughout the experiments is summarised in Table 1.

Predictions, evaluation outputs and aggregate metrics were saved separately for each model version. The base model, first adapter and refinement adapter were also stored as separate files. This prevented later experiments from overwriting earlier results and allowed the models to be evaluated independently.

Table 1: Training, inference and software configuration used in the experiments.

Configuration	Value
Base model	Qwen2.5-1.5B-Instruct
Fine-tuning method	QLoRA
Quantisation	4-bit
Training precision	float16
Training epochs	3
Per-device batch size	2
Random seed	4242564
Original training examples	1,597
Validation examples	204
Test examples	199
Refinement examples	250
GPU	NVIDIA GeForce RTX 4070 (12 GB)
Prolog implementation	SWI-Prolog
Training framework	TRL

The experimental sequence was therefore:

1. evaluate the original Qwen2.5-1.5B-Instruct model;
2. fine-tune the model using QLoRA and the original training split;
3. evaluate the first adapter on the unchanged test set;

4. inspect errors by category and difficulty;
5. create targeted refinement examples;
6. train a second QLoRA adapter using the original and refinement data;
and
7. evaluate the refined adapter using the same test set and metrics.

This procedure provides a consistent comparison of the baseline and initial fine-tuned model using a fixed evaluation pipeline. The refined model was also evaluated on the same test set to measure changes following targeted refinement. However, because the refinement dataset was constructed after analysing error patterns observed on this test set, the refined-model results should not be interpreted as performance on a completely untouched held-out test set. This limitation is considered further when interpreting the results in Chapter 5.

5 Experiments and Results

This chapter presents the experimental results obtained from the baseline model, the first fine-tuned model and the refined model. All three versions were evaluated using the same test set of 199 examples and the metrics described in Chapter 4. Keeping the evaluation data and procedure fixed provides a consistent basis for comparing model outputs. However, as discussed later in this chapter, error patterns observed on the test set informed construction of the refinement dataset, so the final refined-model evaluation is not fully independent of the model-development process.

The experiments were conducted in three stages. First, the original Qwen2.5-1.5B-Instruct model was evaluated without task-specific fine-tuning to establish a baseline. Second, the model was fine-tuned using QLoRA on the original NL-Prolog training set and evaluated under the same conditions. The errors produced by this model were then analysed by difficulty level, category and template group. Finally, a targeted refinement dataset was created from the observed weaknesses, and a second fine-tuned model was trained and evaluated.

5.1 Baseline Evaluation

The first experiment measured how effectively the original Qwen2.5-1.5B-Instruct model could perform NL-Prolog translation without task-specific training. The model received the same instruction used throughout the experiments and was required to return only Prolog code.

The baseline results are shown in Table 2.

Table 2: Baseline evaluation results.	
Metric	Result
Exact Match	1/199 (0.50%)
Syntax Valid	79/199 (39.70%)
Executable	79/199 (39.70%)
Reasoning Correct	6/41 (14.63%)

Only one of the 199 predictions exactly matched the expected Prolog representation. This indicates that the instruction-tuned model had very limited ability to reproduce the dataset’s required mapping between natural-language structures and Prolog without additional training.

Syntax validity was higher than exact-match accuracy. Seventy-nine

outputs, or 39.70%, were accepted as valid Prolog. The same number were executable. The difference between exact match and syntax validity shows that the base model could sometimes generate text that resembled valid Prolog without generating the specific symbolic representation required by the input.

Inspection of individual predictions supported this interpretation. Some outputs contained valid-looking predicates but mapped the natural-language relation incorrectly. For example, a binary relationship could be converted into an unrelated predicate or could use the wrong argument structure. Such outputs may satisfy the Prolog parser while failing to represent the intended meaning.

Reasoning performance was also poor. Only 6 of the 41 Level 4 reasoning examples produced the expected query result, giving an accuracy of 14.63%. This result is important because syntax alone cannot determine whether the formalization represents the original statement correctly. A program may load successfully while encoding incorrect facts or relationships.

The baseline therefore established a clear lower reference point. Qwen2.5-1.5B-Instruct possessed some general ability to generate structured code-like text, but this did not translate into reliable NL-Prolog formalization.

This result is consistent with the motivation behind neuro-symbolic systems such as Logic-LM and LINC. Both report that language models can fail at logical tasks when relied upon directly and instead gain from translating problems into formal representations that external solvers can process. The present baseline additionally demonstrates that generating the intermediate representation is itself a non-trivial task.

5.2 Initial Fine-Tuned Model

The second experiment evaluated the effect of supervised QLoRA fine-tuning. Qwen2.5-1.5B-Instruct was trained for three epochs using the 1,597 examples in the original training split. The resulting adapter was then evaluated on the unchanged 199-example test set.

Performance improved substantially, as shown in Table 3.

Exact-match accuracy increased from 1 correct translation to 162 out of 199. This represents an absolute improvement of 80.91 percentage points.

The improvement in structural correctness was even clearer. Every generated output from the fine-tuned model was accepted as syntactically valid Prolog and was executable. Compared with the baseline, where more than half of the outputs could not be successfully interpreted as Prolog, this sug-

Table 3: Comparison between the baseline and first fine-tuned model.

Metric	Baseline	Fine-Tuned v1
Exact Match	0.50%	81.41%
Syntax Valid	39.70%	100.00%
Executable	39.70%	100.00%
Reasoning Correct	14.63%	87.80%

gests that the model learned the required output syntax reliably from the supervised examples.

Reasoning correctness also increased substantially, from 14.63% to 87.80%. The fine-tuned model therefore improved not only in reproducing the reference strings but also in generating knowledge bases capable of supporting the expected queries.

These results support the use of parameter-efficient fine-tuning for a narrow formalization task. QLoRA allows quantised pretrained language models to be adapted using a comparatively small number of trainable parameters. In this project, the approach allowed a 1.5-billion-parameter model to acquire specialised NL-Prolog generation behaviour using consumer hardware.

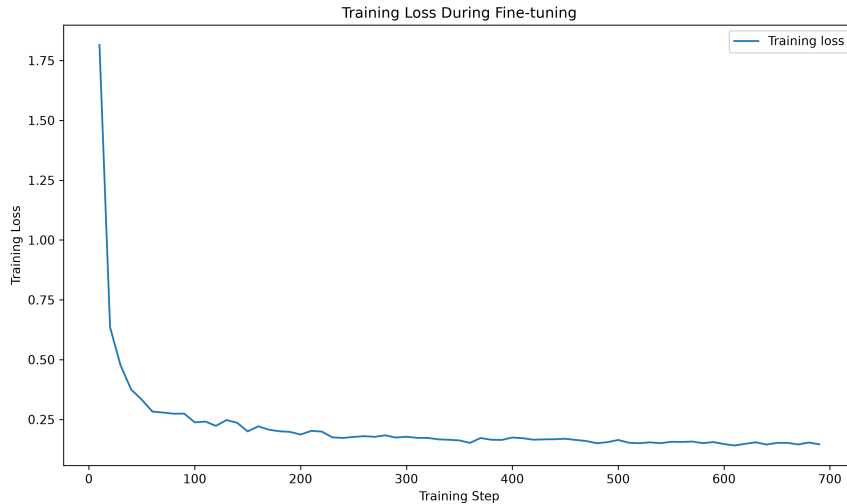


Figure 1: Training loss during QLoRA fine-tuning.

Figure 1 shows the training loss recorded during fine-tuning. Loss de-

creased rapidly during the early training steps before declining more gradually and stabilising towards the end of training. This indicates that most of the initial adaptation occurred early in training, while later steps produced smaller improvements.

5.3 Performance by Difficulty Level

Aggregate accuracy hides substantial variation between translation types. The first fine-tuned model was therefore evaluated separately by category and difficulty level.

Table 4: Exact-match performance of Fine-Tuned v1 by category.

Category	Exact Matches	Accuracy
Fact	43/53	81.13%
Simple Rule	55/60	91.67%
Multi-Condition Rule	23/45	51.11%
Reasoning	41/41	100.00%

Simple rules were translated most reliably among the direct translation categories, reaching 91.67% exact-match accuracy. Facts achieved 81.13%.

The most difficult category was Level 3 multi-condition rules. Only 23 of the 45 examples exactly matched the reference, giving an accuracy of 51.11%. These examples require the model to maintain several variables and conditions simultaneously and may include inequality or negation. An error in any predicate, argument position or condition causes the complete program to fail exact match.

The Level 4 reasoning examples achieved 100% exact match in this category-level analysis, demonstrating that difficulty did not simply increase monotonically with the assigned level. The Level 4 templates were more constrained than the range of Level 3 structures, while Level 3 contained several relational patterns that proved particularly difficult.

This illustrates why reporting only aggregate accuracy would have been insufficient. An overall exact-match score above 80% could suggest relatively uniform performance, while the category analysis shows that the model remained much less reliable on particular symbolic structures.

5.4 Error Analysis

The 37 non-exact predictions from the first fine-tuned model were analysed to determine whether the failures were random or concentrated in particular templates.

The most common failure groups are shown in Table 5.

Table 5: Most common non-exact template groups for Fine-Tuned v1.

Template Group	Non-Exact Predictions
<code>shared_child_inequality</code>	10
<code>employee_not_busy</code>	7
<code>binary_calls</code>	6
<code>bird_not_flightless</code>	5
<code>helps_same_direction</code>	5
<code>binary_admires</code>	4

Some examples contributed to more than one diagnostic grouping, so these counts are useful for identifying recurring patterns rather than as mutually exclusive error classes.

The `shared_child_inequality` template was particularly problematic. These examples require several binary relationships to share variables while also enforcing an inequality condition. Generated programs sometimes reversed the arguments of a relationship or associated a variable with the wrong predicate position. The resulting output could remain syntactically valid while representing a different logical relationship.

Binary predicates were another recurring source of errors. Examples involving `calls/2` and `admires/2` sometimes produced semantically related but incorrect predicate names. For example, a statement requiring `admires` could be represented using a predicate such as `respects`. In ordinary natural-language generation, such a substitution could be considered reasonable. In a formal program, however, `admires/2` and `respects/2` are different predicates unless an explicit relationship between them has been defined.

Similar errors appeared in negative structures. Templates such as `employee_not_busy` and `bird_not_flightless` require the model to preserve negation and the intended predicate simultaneously. A generated program may be valid Prolog but encode a related concept such as `free(X)` instead of the expected logical structure involving `\+ busy(X)`.

These errors demonstrate an important property of the task. Once fine-

tuned, the model rarely failed because it was unable to produce Prolog syntax. Instead, the remaining failures increasingly concerned the mapping between natural-language meaning and exact symbolic structure.

Error analysis plays a similar role in related neuro-symbolic research. Logic-LM uses feedback from the symbolic solver to identify and revise unsuccessful formalizations. The refinement strategy used in this project differs because recurring failures were addressed through additional training examples rather than revising individual outputs during inference.

5.5 Targeted Refinement Dataset

The error analysis motivated a second training stage. Rather than increasing the training set with randomly generated examples, a refinement dataset was constructed to increase exposure to translation patterns related to the observed failures.

The refinement dataset contained 250 validated examples. Its distribution is shown in Table 6.

Table 6: Distribution of the targeted refinement dataset.

Level	Examples
Level 1	90
Level 2	30
Level 3	130
Total	250

The largest proportion was assigned to Level 3 because this was the weakest category in the first evaluation. The refinement set included additional examples involving binary relationships, negation, conjunction and other structures similar to those appearing frequently in the error analysis.

Examples were validated using the same Prolog validation procedure as the main dataset before being included in training. The 250 examples were combined with the original training data, and a second QLoRA adapter was trained. The validation and test sets were left unchanged.

This design allowed the refinement experiment to test a specific question: whether additional examples concentrated on known failure patterns could improve performance without modifying the evaluation set.

5.6 Refined Model Results

The refined model, referred to as **Fine-Tuned v2**, was evaluated using the same 199 test examples.

The results are shown in Table 7.

Table 7: Comparison between Fine-Tuned v1 and Fine-Tuned v2.

Metric	Fine-Tuned v1	Fine-Tuned v2
Exact Match	81.41%	87.94%
Syntax Valid	100.00%	100.00%
Executable	100.00%	100.00%
Reasoning Correct	87.80%	100.00%

Exact-match accuracy increased from 162 to 175 correct predictions. This corresponds to an improvement of 6.53 percentage points over the first fine-tuned model and 87.44 percentage points over the original baseline.

All 199 outputs remained syntactically valid and executable. The additional fine-tuning therefore improved translation accuracy without reducing the structural reliability achieved by the first model.

Most notably, reasoning correctness increased to 41 out of 41 examples. All Level 4 queries therefore produced their expected results under the final model.

These final results require an important qualification. The refinement dataset was constructed after analysing recurring error patterns observed on the original test set. Although the individual test examples were not directly included in the refinement training data, information about the structures and categories on which the first fine-tuned model performed poorly influenced the second training stage. The Fine-Tuned v2 results should therefore not be interpreted as performance on a completely untouched held-out benchmark.

The category-level exact-match results are shown in Table 8.

Simple-rule performance increased from 91.67% to 100%, while fact translation increased from 81.13% to 86.79%.

Level 3 remained the weakest category, but improved from 51.11% to 62.22%. This is an increase of 11.11 percentage points and represents five additional exact translations on the 45 Level 3 test examples.

Figure 2 provides a direct comparison of exact-match accuracy by category before and after refinement. The largest improvement occurred for multi-condition rules, although this category continued to perform substan-

Table 8: Exact-match performance of Fine-Tuned v2 by category.

Category	Exact Matches	Accuracy
Fact	46/53	86.79%
Simple Rule	60/60	100.00%
Multi-Condition Rule	28/45	62.22%
Reasoning	41/41	100.00%

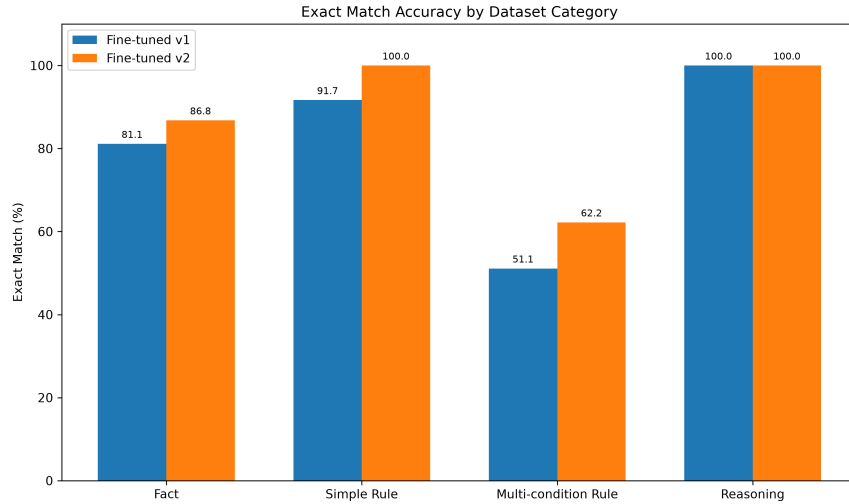


Figure 2: Exact-match accuracy by dataset category for Fine-Tuned v1 and Fine-Tuned v2.

tially below the others. Simple rules reached 100%, while reasoning examples remained at 100% across both fine-tuned versions. The number of non-exact outputs across the whole test set fell from 37 to 24.

5.7 Remaining Errors

Although refinement improved overall performance, several failure patterns remained.

The most common final failure groups are shown in Table 9.

The continued presence of all 10 `shared_child_inequality` failures is particularly notable. Unlike several simpler categories, targeted additional training did not resolve this structure on the test examples. The pattern combines multiple binary predicate applications, shared variables and in-

Table 9: Most common remaining failure groups for Fine-Tuned v2.

Template Group	Remaining Failures
shared_child_inequality	10
binary_admires	4
bird_not_flightless	4
employee_not_busy	3
binary_calls	3

equality, making errors in variable binding or argument direction especially consequential.

By contrast, some categories responded strongly to refinement. Simple rules reached complete exact-match accuracy, and the number of errors involving several other templates declined. This suggests that targeted additional examples can correct some systematic translation behaviours, although different structures respond differently to additional supervision.

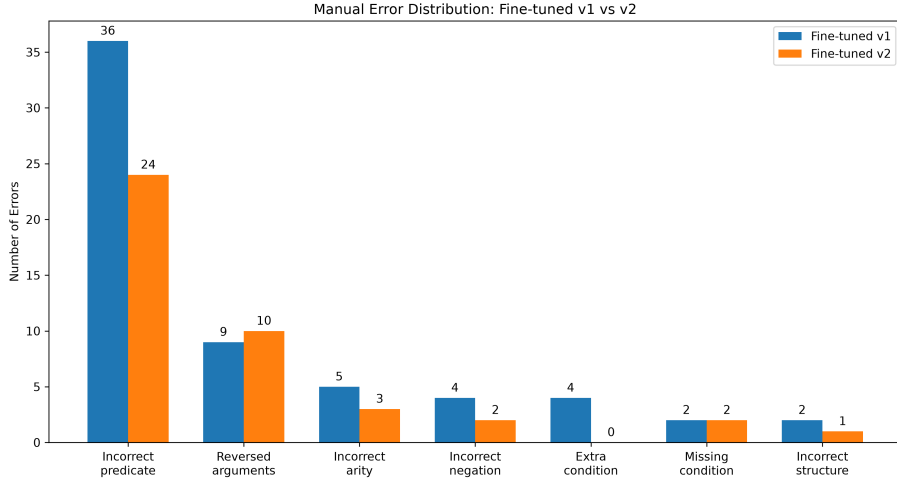


Figure 3: Distribution of manually identified errors for Fine-Tuned v1 and Fine-Tuned v2.

Figure 3 provides a broader comparison of the manually identified error types before and after refinement. Incorrect predicate selection remained the most common error but decreased substantially in the refined model. Incorrect arity, negation, additional conditions and incorrect structures also became less frequent. Reversed arguments remained problematic and in-

creased slightly from nine to ten identified cases. A single prediction could contain more than one error type, so the counts in the figure do not sum to the total number of non-exact predictions.

The final 24 non-exact examples also reinforce the distinction between formal validity and exact semantic translation. Despite these failures, every output remained syntactically valid and executable. The remaining challenge was therefore not primarily producing legal Prolog syntax. It was producing the precise predicates, argument relationships and logical operations specified by the natural-language input.

5.8 Overall Comparison

The complete experimental progression is summarised in Table 10.

Table 10: Overall comparison of the three evaluated model versions.

Model	Exact Match	Syntax Valid	Executable	Reasoning Correct
Baseline	0.50%	39.70%	39.70%	14.63%
Fine-Tuned v1	81.41%	100.00%	100.00%	87.80%
Fine-Tuned v2	87.94%	100.00%	100.00%	100.00%

The largest improvement occurred between the baseline and first fine-tuned model. This indicates that supervised task-specific adaptation was responsible for most of the performance gain. Refinement produced a smaller but still measurable additional improvement, particularly for simple and multi-condition rules.

Figure 4 visualises the overall progression across the three model versions. Fine-tuning produced a substantial improvement across all four evaluation metrics, particularly when compared with the original baseline. The refinement stage produced a further increase in exact-match and reasoning accuracy while maintaining 100% syntax validity and executability.

The results also provide evidence for the usefulness of evaluating symbolic generation at more than one level. If only syntax validity had been measured, both fine-tuned models would appear identical at 100%. Exact-match evaluation reveals substantial differences between them. Conversely, exact match alone would not show that all final generations formed executable Prolog programs or that every reasoning query produced its expected result.

A related motivation appears in LINC, where language models are used as semantic parsers and external theorem provers perform the deductive

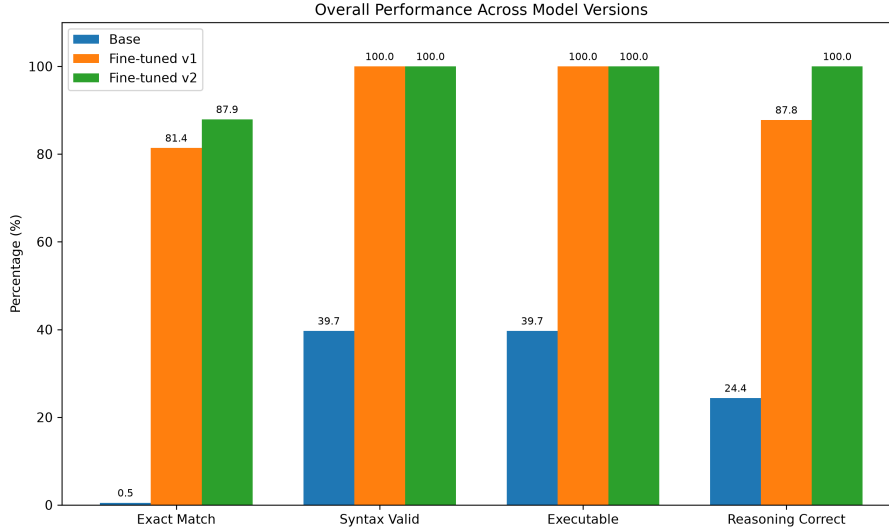


Figure 4: Overall performance of the baseline, Fine-Tuned v1 and Fine-Tuned v2 models.

reasoning. This modular approach demonstrates how symbolic execution can complement neural language processing, provided that the intermediate formalization is sufficiently accurate.

The experiments in this project examine the preceding problem more directly: whether a small language model can be trained to produce that intermediate symbolic representation reliably. The progression from 0.50% to 87.94% exact match, together with complete syntactic validity and executability in the final model, shows that Qwen2.5-1.5B can acquire strong NL-Prolog translation behaviour from a relatively small, controlled dataset. At the same time, the remaining Level 3 failures show that syntactically correct generation does not guarantee precise semantic formalization. The implications and limitations of these findings are considered in Chapter 6.

6 Discussion

The experiments demonstrate that a small language model can be adapted to perform natural-language-to-Prolog translation with substantially greater reliability than the original instruction-tuned model. Fine-tuning Qwen2.5-1.5B-Instruct increased exact-match accuracy from 0.50% to 81.41%, while targeted refinement increased this further to 87.94%. Both fine-tuned models also generated syntactically valid and executable Prolog for every example in the test set. These results answer the central research question positively within the controlled setting of this project: a 1.5-billion-parameter language model can learn a strong mapping from natural-language reasoning statements to executable Prolog through task-specific training.

However, the different evaluation metrics and remaining errors show that this conclusion requires qualification. The model learned Prolog syntax more completely than it learned the exact semantic mappings represented in the dataset. The results also depend on a synthetic and relatively constrained dataset, and the refinement procedure introduces limitations in how the final test performance should be interpreted.

6.1 Effect of Fine-Tuning

The largest improvement occurred between the baseline and the first fine-tuned model. Exact match increased by more than 80 percentage points, while syntax validity and executability both increased from 39.70% to 100%. This suggests that task-specific supervision was much more important than the model’s existing general ability to generate code.

The baseline demonstrates that instruction tuning alone was insufficient for this specialised task. Qwen2.5-1.5B-Instruct occasionally generated valid Prolog-like output, but it rarely produced the expected formal representation. After being exposed to 1,597 paired examples, the model consistently followed the required syntax and usually mapped the natural-language input to the expected predicates and rules.

This supports the broader principle behind parameter-efficient adaptation. QLoRA was originally proposed as a way of adapting quantised language models while reducing the memory required for training. In this project, it allowed task-specific behaviour to be learned using a consumer GPU rather than requiring full-parameter training. The result is significant for the aim of the project because it suggests that useful autoformalization behaviour is not restricted to very large models or high-resource training environments.

The result should not, however, be interpreted as showing that a 1.5B model has acquired general natural-language-to-Prolog translation ability. The training and evaluation language is deliberately constrained. The stronger conclusion supported by the experiments is that a small model can learn this mapping effectively within a controlled domain.

6.2 Symbolic Execution as an Evaluation Tool

One of the clearest findings is the value of evaluating generated formalizations at more than one level. Exact match, syntax validity, executability and reasoning correctness measure related but different properties.

The final model achieved 87.94% exact match while achieving 100% syntax validity and executability. Therefore, 24 outputs differed from the reference despite still forming legal Prolog programs. This distinction would have been hidden if only a single metric had been used.

Execution-based evaluation is particularly appropriate for formal outputs because it takes advantage of information that is unavailable when evaluating unrestricted natural language. Previous work on natural-language-to-code translation has similarly demonstrated that program execution can provide useful semantic information beyond execution-unaware comparison. Candidate programs can, for example, be distinguished according to their execution behaviour rather than assuming that surface similarity alone determines correctness.

The findings of this project support the same general principle. An exact textual comparison is useful because the dataset defines a specific intended representation, but it is strict. Two programs may differ textually while remaining executable. At the same time, executability alone is too weak: a program can execute successfully while representing the wrong relationship.

The remaining errors demonstrate this directly. Generating `respects(X,Y)` instead of `admires(X,Y)` may produce perfectly valid Prolog, yet the symbolic meaning differs from the reference. Similarly, reversing arguments in `parent/2` produces a syntactically valid program representing a different relationship. Therefore, execution should complement rather than replace semantic evaluation.

6.3 Reasoning Correctness

Reasoning accuracy showed one of the largest improvements, increasing from 14.63% for the baseline to 87.80% for the first fine-tuned model and 100% for the refined model. This provides evidence that the generated programs

were not merely syntactically well formed. For the 41 reasoning examples in the final test set, every program supported the expected query result.

This result is consistent with the motivation of neuro-symbolic approaches such as Logic-LM and LINC. Logic-LM separates translation from inference by having a language model create a symbolic formulation before a deterministic solver performs reasoning. LINC similarly treats the language model as a semantic parser and delegates deduction to a first-order logic prover.

The present system follows the same broad separation but focuses on improving the translation stage through fine-tuning. Once a correct Prolog representation has been generated, SWI-Prolog can perform the inference deterministically.

Nevertheless, 100% reasoning accuracy should not be interpreted as evidence that the model can solve arbitrary logical problems. Only 41 reasoning examples from the test set were evaluated, and they originated from the controlled Level 4 templates used by the project. The result demonstrates complete performance on this test subset rather than universal reasoning capability.

6.4 Remaining Difficulties

The distribution of errors reveals that formalization difficulty depends strongly on logical structure. Simple rules achieved 100% exact match in the refined model, while multi-condition rules reached only 62.22%.

The persistent `shared_child_inequality` failures are particularly informative. These examples require the model to maintain several predicates, variable relationships and an inequality constraint simultaneously. All ten examples in this failure group remained incorrect after refinement. This suggests that adding more examples alone did not completely solve the underlying structural problem.

Errors involving binary predicates reveal another challenge. Natural language allows concepts such as “admires” and “respects” to be treated as semantically related. Formal programming languages do not provide this flexibility unless that relationship has been explicitly defined. Autoformalization therefore requires a model to suppress some of its normal linguistic behaviour and preserve exact predicate identity.

This supports an observation made in related neuro-symbolic work: the semantic-parsing stage remains a major failure point even when a symbolic prover can perform deduction reliably. The symbolic component removes uncertainty from the inference stage, but it cannot correct a program that

encodes the wrong meaning.

Recent work using Prolog provides a further comparison. Research involving language models using Prolog as an external reasoning tool has highlighted a distinction between answer accuracy and the production of fully auditable symbolic programs. Obtaining a correct answer does not necessarily mean that the intermediate symbolic representation is itself faithful. This reinforces the decision in the present project to evaluate the generated Prolog directly rather than judging only the final reasoning result.

6.5 Effect of Targeted Refinement

The refinement experiment increased exact-match accuracy from 81.41% to 87.94%, reduced non-exact outputs from 37 to 24 and increased reasoning correctness to 100%. The improvement indicates that error analysis can be used to guide additional training more efficiently than simply generating more examples without considering model behaviour.

This is conceptually related to approaches that use symbolic-solver feedback to improve failed formalizations. The mechanism used here differs: instead of revising a failed output during inference, recurring errors were identified across the evaluation results and used to construct new supervised training examples.

However, this refinement procedure also creates an important limitation. The errors used to decide which examples required additional training were observed on the same test set later used to report the Fine-Tuned v2 results. Although the exact test examples themselves were not added to the training set, information about their failure categories influenced the refinement data. The v2 test set was therefore no longer completely independent of model development.

Consequently, the 87.94% final exact-match result should be interpreted primarily as evidence that targeted refinement improved performance on the identified weaknesses. A stronger estimate of final generalisation would require a second untouched test set that was not inspected until all refinement decisions had been completed. Future work should use the validation or development set for error-driven refinement and reserve an independent final test set solely for the last evaluation.

6.6 Dataset and Experimental Limitations

The primary limitation of the project is the controlled nature of the dataset. Synthetic generation made it possible to know the intended Prolog represen-

tation, automatically validate every target and systematically vary logical structures. However, the resulting natural language is less diverse than unrestricted human writing.

The model may therefore have learned recurring correspondences between templates and Prolog structures rather than developing a general ability to formalize arbitrary natural language. Group-aware splitting reduces direct leakage between related template examples, but it cannot completely reproduce the variation found in naturally written reasoning problems.

The test set is also relatively small at 199 examples. Category-level results are therefore based on much smaller subsets, including only 45 multi-condition rules and 41 reasoning examples. Additional evaluation on independently written examples, external datasets or deliberately adversarial paraphrases would provide stronger evidence about generalisation.

A further limitation is the loss of test-set independence during targeted refinement. As discussed above, the test examples themselves were not added to the refinement dataset, but their observed failure patterns influenced its construction. The refined-model results may therefore provide an optimistic estimate of performance compared with evaluation on a completely unseen dataset.

Only one base model size and one main training configuration were also evaluated. The experiments therefore cannot determine whether Qwen2.5-1.5B-Instruct is optimal, how performance scales with model size, or whether different LoRA parameters, training durations or dataset sizes would produce better results. These questions were outside the practical scope of the project but provide clear directions for further work.

Overall, the experiments demonstrate that fine-tuning can make a small language model highly reliable at generating executable Prolog within a controlled domain. They also show why symbolic validation is valuable: Prolog execution makes structural failures directly observable and supports functional reasoning tests. The remaining errors, particularly in multi-condition relations, indicate that the more difficult problem is not learning Prolog syntax but preserving precise semantics when translating increasingly complex natural-language structures.

7 Conclusion and Future Work

This project investigated whether a small language model could be trained to translate natural-language reasoning statements into executable Prolog. The motivation was to explore autoformalization as a way of combining the natural-language capabilities of language models with the explicit and verifiable reasoning provided by symbolic systems. Rather than relying on a language model to perform both interpretation and reasoning internally, the developed system uses the model to generate a formal Prolog representation that can subsequently be checked and executed using SWI-Prolog.

A controlled dataset of 2,000 natural-language and Prolog pairs was constructed for this purpose. The examples were divided into four levels covering facts, simple rules, multi-condition rules and reasoning problems. Automated validation with SWI-Prolog ensured that all target programs were syntactically valid and executable, while reasoning examples were additionally checked against their expected query results. A group-aware training, validation and test split was used to reduce direct leakage between closely related generated templates.

The experiments used Qwen2.5-1.5B-Instruct as the base model. Before fine-tuning, its performance on the 199-example test set was poor. Exact-match accuracy was only 0.50%, while 39.70% of outputs were syntactically valid and executable. Reasoning correctness was 14.63%. These results demonstrated that general instruction-following ability alone was insufficient for reliable natural-language-to-Prolog translation.

QLoRA fine-tuning produced a substantial improvement. The first fine-tuned model achieved 81.41% exact-match accuracy, while every generated program was syntactically valid and executable. Reasoning accuracy increased to 87.80%. Error analysis showed that the remaining failures were concentrated in particular structures rather than being evenly distributed across the dataset. Multi-condition rules were especially difficult, with recurring errors involving binary predicate direction, inequality and negation.

These observations were used to construct an additional refinement dataset containing 250 targeted examples. Following a second training stage, exact-match accuracy increased to 87.94%. Syntax validity and executability remained at 100%, while all 41 reasoning examples produced their expected results. Simple-rule exact match reached 100%, although multi-condition rules remained the weakest category at 62.22%. This suggests that the model learned the basic syntax and structure of Prolog more reliably than some of the more complex semantic relationships expressed within it.

The results therefore provide a positive answer to the central research

question within the scope of the experiment. A 1.5-billion-parameter language model can be fine-tuned to perform NL-Prolog autoformalization with high reliability on a controlled domain, and parameter-efficient training makes this possible using consumer hardware. The project also demonstrates the usefulness of Prolog execution as an evaluation mechanism. Exact match alone does not fully describe model performance: an output can differ from the reference while remaining executable, while an executable program may still encode an incorrect semantic relationship. Combining exact match with syntax, execution and reasoning evaluation therefore gives a more informative picture of autoformalization quality.

There are several limitations to these conclusions. The dataset was synthetically generated from controlled templates and consequently does not represent the full linguistic variation of naturally written reasoning problems. The test set was also relatively small. In addition, the refinement examples were designed after inspecting weaknesses observed on the test set, meaning that the final 87.94% result should not be interpreted as a fully independent estimate of generalisation. The refinement experiment demonstrates improvement on identified weaknesses, but a completely untouched final test set would be required for stronger evaluation.

Future work should therefore evaluate the model on natural or independently written language, including paraphrases and structures not represented by the original templates. A stronger experimental design could use a development set for error-driven refinement while reserving a second test set until all model development is complete. Further experiments could also compare different model sizes, training-set sizes and parameter-efficient fine-tuning configurations.

The most persistent Level 3 errors provide another direction for future research. Additional work could investigate more structured training for variable bindings and argument ordering, or use feedback from SWI-Prolog during generation and refinement rather than only after prediction. A system could, for example, detect invalid or unsuccessful generated programs and request a corrected formalization automatically. This would extend the current pipeline from execution-based evaluation towards execution-guided generation.

Overall, the project demonstrates that relatively small language models can learn a useful interface between natural language and symbolic logic. Although the resulting model is not a general-purpose autoformalizer, the substantial improvement from the untuned baseline and the reliability of executable output show that combining task-specific language-model training with symbolic validation is a practical approach to producing more checkable

forms of machine reasoning.

References

- [1] Qiming Bao, Alex Yuxuan Peng, Tim Hartill, Neset Tan, Zhenyun Deng, Michael Witbrock, and Jiamou Liu. Multi-step deductive reasoning over natural language: An empirical study on out-of-distribution generalisation. In *Proceedings of the 16th International Workshop on Neural-Symbolic Learning and Reasoning*, pages 202–217, 2022.
- [2] Peter Clark, Oyvind Tafjord, and Kyle Richardson. Transformers as soft reasoners over language. In *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence (IJCAI-20)*, pages 3882–3890, 2020.
- [3] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. Qlora: Efficient finetuning of quantized llms. In *Advances in Neural Information Processing Systems*, volume 36, pages 10088–10115, 2023.
- [4] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- [5] Niklas Mellgren, Peter Schneider-Kamp, and Lukas Galke Poech. Training language models to use prolog as a tool. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 32860–32880, 2026.
- [6] Theo X. Olausson, Alex Gu, Benjamin Lipkin, Cedegao E. Zhang, Armando Solar-Lezama, Joshua B. Tenenbaum, and Roger Levy. Linc: A neurosymbolic approach for logical reasoning by combining language models with first-order logic provers. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5153–5176, 2023.
- [7] Liangming Pan, Alon Albalak, Xinyi Wang, and William Yang Wang. Logic-lm: Empowering large language models with symbolic solvers for faithful logical reasoning. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 3806–3824, 2023.

- [8] Qwen Team. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- [9] Jan Wielemaker, Tom Schrijvers, Markus Triska, and Torbjorn Lager. Swi-prolog. *Theory and Practice of Logic Programming*, 12(1–2):67–96, 2012.
- [10] Yuhuai Wu, Albert Qiaochu Jiang, Wenda Li, Markus N. Rabe, Charles Staats, Mateja Jamnik, and Christian Szegedy. Autoformalization with large language models. In *Advances in Neural Information Processing Systems*, volume 35, pages 32353–32368, 2022.
- [11] Sen Yang, Xin Li, Leyang Cui, Lidong Bing, and Wai Lam. Neuro-symbolic integration brings causal and reliable reasoning proofs. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 5747–5759, 2025.